

Polynomial Regression

Tian Han

Outline

- Polynomial Regression
- Binary Classification - Logistic regression

Warm-up: Linear Regression


Linear Regression (Task)

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$

Output: a vector $\mathbf{w} \in \mathbb{R}^d$ and scalar $b \in \mathbb{R}$ such that $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$.

Tasks

assume y_i is a linear function of \mathbf{x}_i .



Linear
Regression

Least Squares Regression (Method)

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$

1. Add one dimension to $\mathbf{x} \in \mathbb{R}^d$: $\bar{\mathbf{x}}_j = [\mathbf{x}_j; 1] \in \mathbb{R}^{d+1}$.
2. Solve least squares regression: $\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \mathbf{w} - \mathbf{y}\|_2^2$.

Tasks

Methods

Linear
Regression

Least Squares Regression

Least Squares Regression (Method)

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$

1. Add one dimension to $\mathbf{x} \in \mathbb{R}^d$: $\bar{\mathbf{x}}_j = [\mathbf{x}_j; 1] \in \mathbb{R}^{d+1}$.
2. Solve least squares regression: $\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \mathbf{w} - \mathbf{y}\|_2^2$.

Tasks

Linear
Regression

Methods

Least Squares Regression

Algorithms

Analytical Solution

Gradient Descent

Polynomial Regression

The Regression Task

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$.

Output: a function $f: \mathbb{R}^d \mapsto \mathbb{R}$ such that $f(\mathbf{x}) \approx y$.

Question: f is unknown! So how to learn f ?

The Regression Task

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$.

Output: a function $f: \mathbb{R}^d \mapsto \mathbb{R}$ such that $f(\mathbf{x}) \approx y$.

Question: f is unknown! So how to learn f ?

Answer: polynomial approximation; f is a polynomial function.

Taylor expansion: $f(x) = f(a) + f'(a)(a - x) + \frac{f''(a)}{2!}(a - x)^2 + \dots$

Polynomial Regression: 1D Example

Input: scalars $x_1, \dots, x_n \in \mathbb{R}$ and labels $y_1, \dots, y_n \in \mathbb{R}$.

Output: a function $f: \mathbb{R} \mapsto \mathbb{R}$ such that $f(x) \approx y$.

One-dimensional example: $f(x) = w_0 + w_1x + w_2x^2 + \dots + w_px^p$.

Polynomial Regression: 1D Example

Input: scalars $x_1, \dots, x_n \in \mathbb{R}$ and labels $y_1, \dots, y_n \in \mathbb{R}$.

Output: a function $f: \mathbb{R} \mapsto \mathbb{R}$ such that $f(x) \approx y$.

One-dimensional example: $f(x) = w_0 + w_1x + w_2x^2 + \dots + w_px^p$.

Polynomial regression:

1. Define a feature map $\boldsymbol{\phi}(x) = [1, x, x^2, x^3, \dots, x^p]$.
2. For $j = 1$ to n , do the mapping $x_j \mapsto \boldsymbol{\phi}(x_j)$.
 - Let $\boldsymbol{\Phi} = [\boldsymbol{\phi}(x_1); \dots, \boldsymbol{\phi}(x_n)]^T \in \mathbb{R}^{n \times (p+1)}$
3. Solve the least squares regression $\min_{\mathbf{w} \in \mathbb{R}^{p+1}} \|\boldsymbol{\Phi} \mathbf{w} - \mathbf{y}\|_2^2$.

Polynomial Regression: 2D Example

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^2$ and labels $y_1, \dots, y_n \in \mathbb{R}$.

Output: a function $f: \mathbb{R}^2 \mapsto \mathbb{R}$ such that $f(\mathbf{x}_i) \approx y_i$.

Two-dimensional example: how to do feature mapping?

Polynomial features:

$$\Phi(\mathbf{x}) = [1, \underbrace{x_1, x_2}_{\text{degree-1}}, \underbrace{x_1^2, x_2^2, x_1 x_2}_{\text{degree-2}}, \underbrace{x_1^3, x_2^3, x_1 x_2^2, x_1^2 x_2}_{\text{degree-3}}].$$

degree-0 degree-1 degree-2 degree-3

Polynomial Regression

```
import numpy
X = numpy.arange(6).reshape(3, 2)
print('X = ')
print(X)
```

```
X =
[[0 1]
 [2 3]
 [4 5]]
```

Polynomial Regression

```
import numpy
X = numpy.arange(6).reshape(3, 2)
print('X = ')
print(X)
```

```
X =
[[0 1]
 [2 3]
 [4 5]]
```

```
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=3)
Phi = poly.fit_transform(X)
print('Phi = ')
print(Phi)
```

```
Phi =
[[ 1.  0.  1.  0.  0.  1.  0.  0.  0.  1.]
 [ 1.  2.  3.  4.  6.  9.  8. 12. 18. 27.]
 [ 1.  4.  5. 16. 20. 25. 64. 80. 100. 125.]]
```

degree-0 degree-1 degree-2 degree-3


Polynomial Regression

- \mathbf{x} : d -dimensional
- $\Phi(\mathbf{x})$: degree- p polynomial
- The dimension of $\Phi(\mathbf{x})$ is $O(d^p)$

Polynomial Regression

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$.

Output: a function $f: \mathbb{R}^d \mapsto \mathbb{R}$ such that $f(\mathbf{x}_i) \approx y_i$.



Training, Test, and Overfitting

Polynomial Regression: Training

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$.

Feature map: $\phi(\mathbf{x})$. Its dimension is $O(d^p)$.

Least squares: $\min_{\mathbf{w}} \|\Phi \mathbf{w} - \mathbf{y}\|_2^2$.

Polynomial Regression: Training

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$.


Feature map: $\phi(\mathbf{x})$. Its dimension is $O(d^p)$.

Least squares: $\min_{\mathbf{w}} \|\Phi \mathbf{w} - \mathbf{y}\|_2^2$.



Question: what will happen as p grows?

1. For sufficiently large p , the dimension of the feature $\phi(\mathbf{x})$ exceeds n .
2. Then you can find \mathbf{w} such that $\Phi \mathbf{w} = \mathbf{y}$. (Zero training error!)

Training and Testing



Underfitting



Overfitting

Training and Testing

Train:

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$.


Output: a function $f: \mathbb{R}^d \mapsto \mathbb{R}$ such that $f(\mathbf{x}_i) \approx y_i$.

Test:


Input: a *never-seen-before* feature vectors $\mathbf{x}' \in \mathbb{R}^d$.

output: predict its label by $f(\mathbf{x}')$.


Training and Testing



BAD




GOOD



BAD


Training and Testing



Hyper-Parameter Tuning

Question: for the polynomial regression model, how to determine the degree p ?

Answer: the degree p leads to the smallest test error.



Hyper-Parameter Tuning

Training Set

Test Set

Train a degree-1 polynomial regression



Test MSE = 23.2

Train a degree-2 polynomial regression



Test MSE = 19.0

Train a degree-3 polynomial regression



Test MSE = 16.7

Train a degree-4 polynomial regression



Test MSE = 12.2

Train a degree-5 polynomial regression



Test MSE = 14.8

Train a degree-6 polynomial regression



Test MSE = 25.1

Train a degree-7 polynomial regression



Test MSE = 39.4

Train a degree-8 polynomial regression



Test MSE = 53.0

Hyper-Parameter Tuning

Training Set

Test Set

Train a degree-1 polynomial regression



Test MSE = 23.2

Train a degree-2 polynomial regression



Test MSE = 19.0

Train a degree-3 polynomial regression



Test MSE = 16.7

Train a degree-4 polynomial regression



Test MSE = 12.2

Train a degree-5 polynomial regression



Test MSE = 14.8

Train a degree-6 polynomial regression



Test MSE = 25.1

Train a degree-7 polynomial regression



Test MSE = 39.4

Train a degree-8 polynomial regression



Test MSE = 53.0

Hyper-Parameter Tuning

Training Set


Test Set

Train a degree-1 polynomial regression



Test MSE = 23.2

Train a degree-2 polynomial regression



Test MSE = 19.0

Train a degree-3 polynomial regression



Test MSE = 16.7

Train a degree-4 polynomial regression



Test MSE = 12.2

Train a degree-5 polynomial regression




Test MSE = 14.8

Train a degree-6 polynomial regression




Test MSE = 25.1

Train a degree-7 polynomial regression



Test MSE = 39.4

Train a degree-8 polynomial regression



Test MSE = 53.0

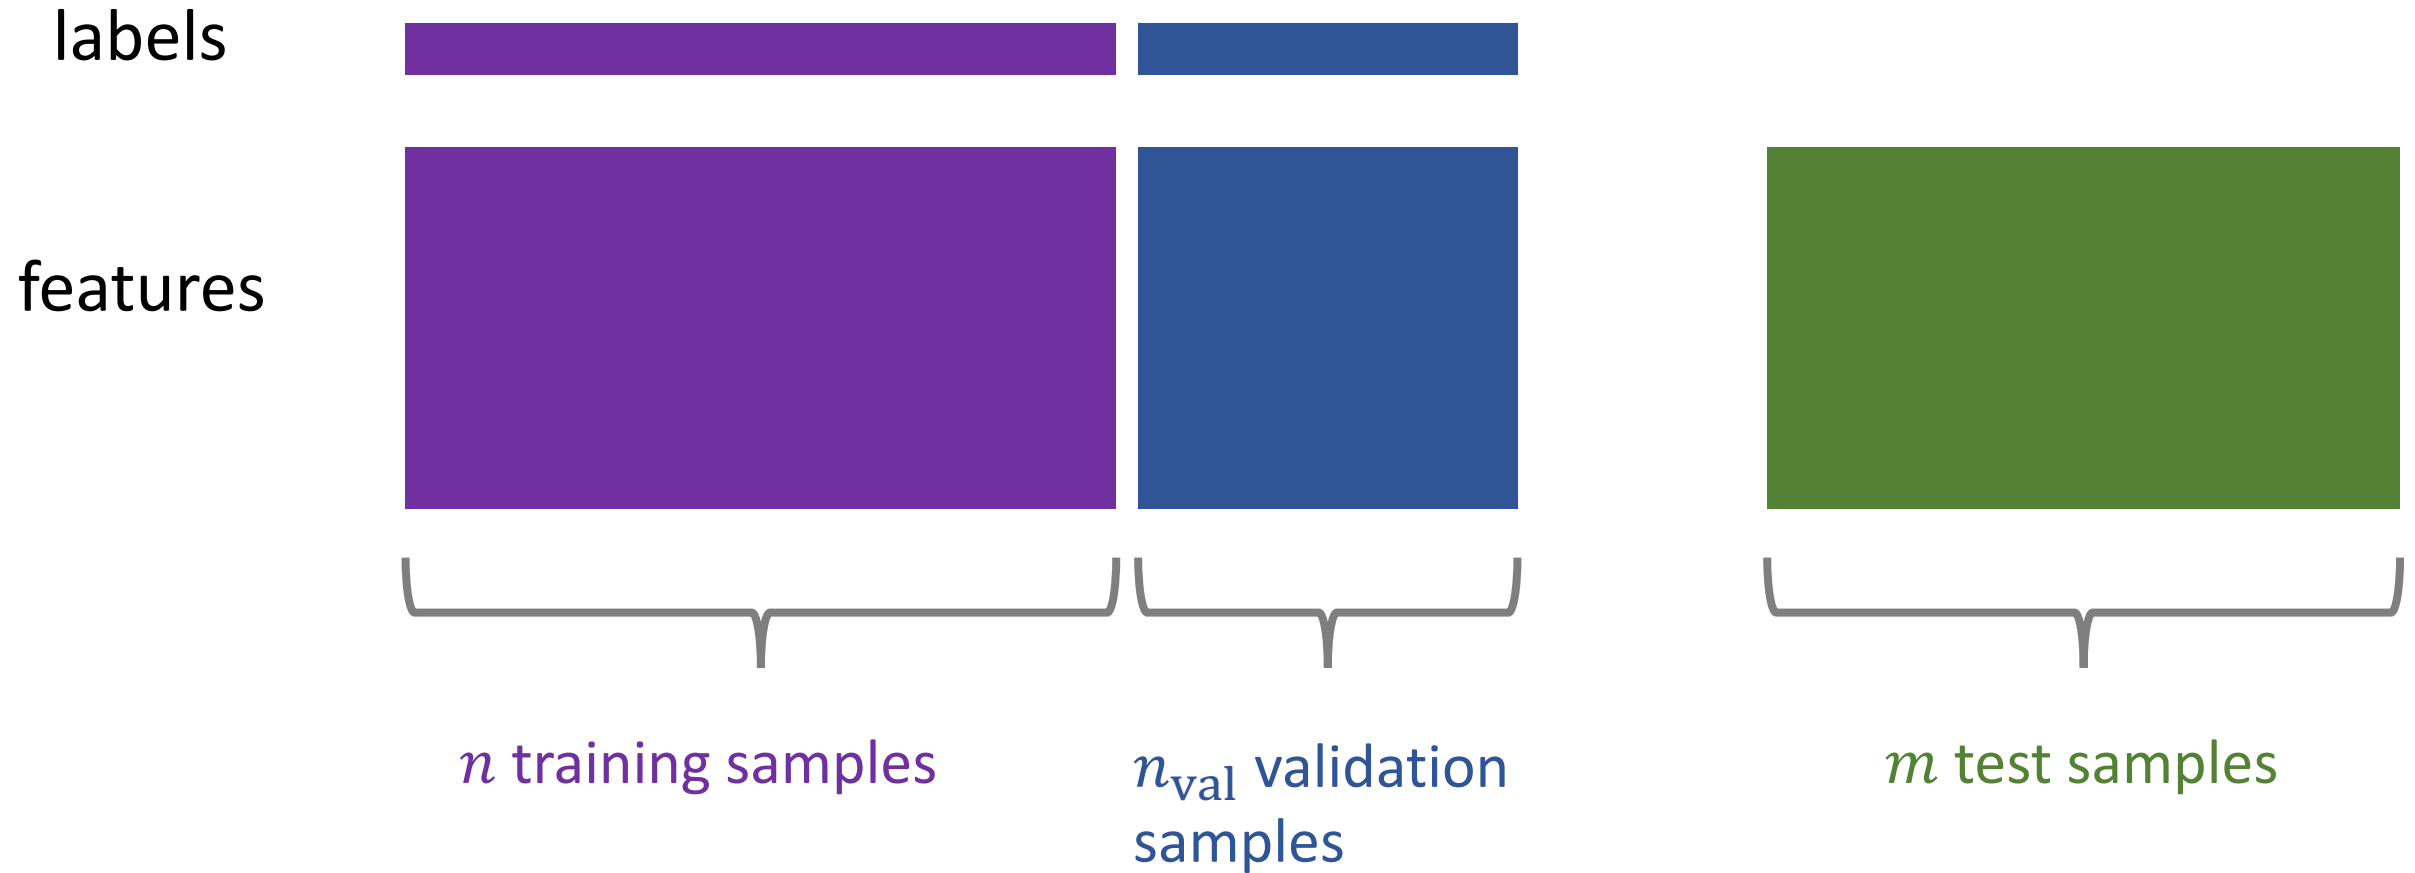
- **Wrong! The test labels are unavailable!**
- **Even if you have the test labels, never do this!**

Cross-Validation (Naïve Approach) for Hyper-Parameter Tuning

Cross-Validation (Naïve Approach)



Cross-Validation (Naïve Approach)



Cross-Validation (Naïve Approach)

Training Set

Test~~Set~~

Train a degree-1 polynomial regression



Test MSE~~X~~ = 23.2

Train a degree-2 polynomial regression



Test MSE~~X~~ = 19.0

Train a degree-3 polynomial regression



Test MSE~~X~~ = 16.7

Train a degree-4 polynomial regression



Test MSE~~X~~ = 12.2

Train a degree-5 polynomial regression



Test MSE~~X~~ = 14.8

Train a degree-6 polynomial regression



Test MSE~~X~~ = 25.1

Train a degree-7 polynomial regression



Test MSE~~X~~ = 39.4

Train a degree-8 polynomial regression



Test MSE~~X~~ = 53.0

Cross-Validation (Naïve Approach)

Training Set

Validation Set

~~Test Set~~

Train a degree-1 polynomial regression



Valid. MSE = 23.1

Train a degree-2 polynomial regression



Valid. MSE = 19.2

Train a degree-3 polynomial regression



Valid. MSE = 16.3

Train a degree-4 polynomial regression



Valid. MSE = 12.5

Train a degree-5 polynomial regression



Valid. MSE = 14.4

Train a degree-6 polynomial regression



Valid. MSE = 25.0

Train a degree-7 polynomial regression



Valid. MSE = 39.1

Train a degree-8 polynomial regression

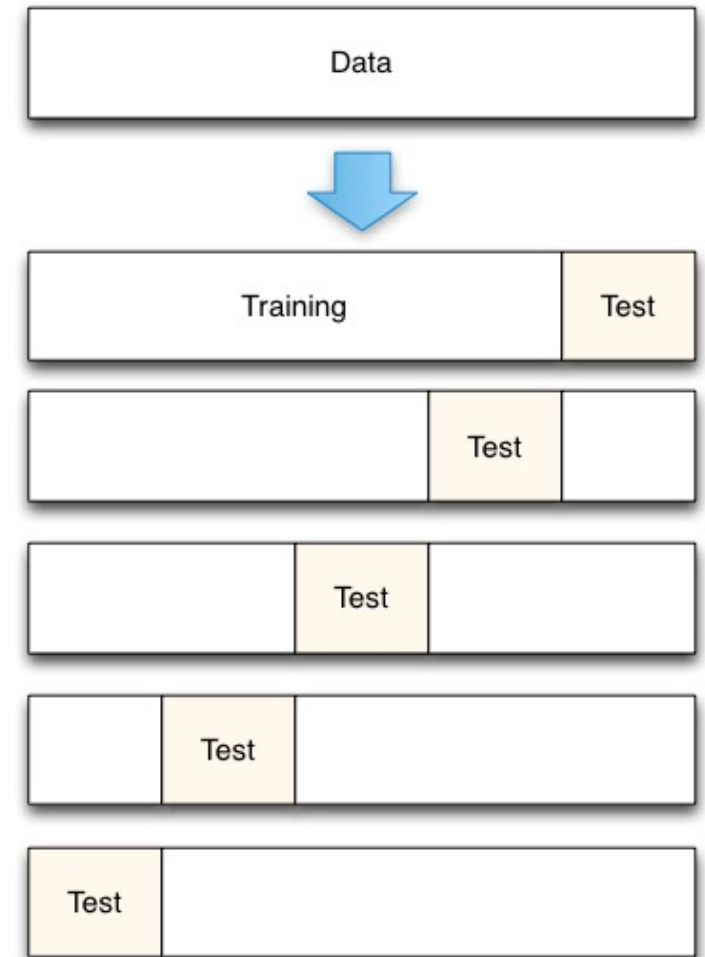


Valid. MSE = 53.5

***k*-Fold Cross-Validation**

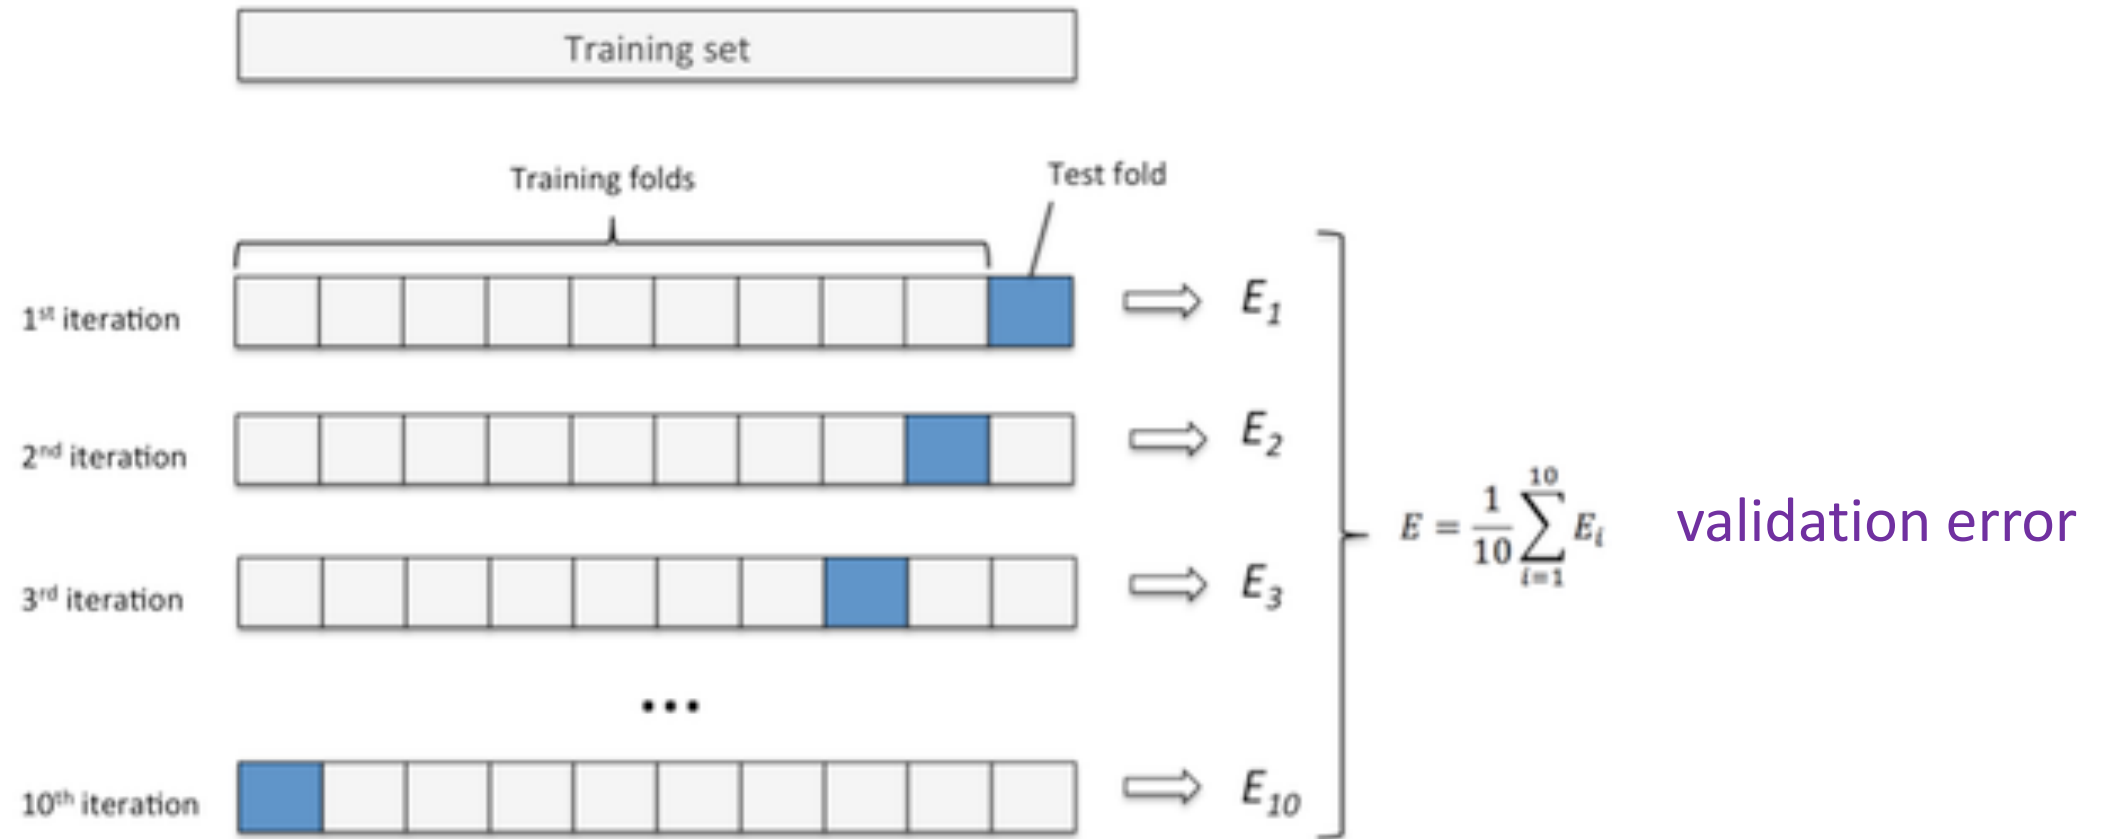
k -Fold Cross-Validation

1. Propose a grid of hyper-parameters.
 - E.g. $p \in \{1, 2, 3, 4, 5, 6\}$.
2. Randomly partition the training samples to k parts.
 - $k - 1$ parts for training.
 - One part for test.
3. Compute the averaged test errors of the k repeats.
 - The average is called the **validation error**.
4. Choose the hyper-parameter p that leads to the smallest **validation error**.



Example: 5-fold cross-validation

Example: 10-Fold Cross-Validation



Example: 10-Fold Cross-Validation

hyper-parameter	validation error
-----------------	------------------

p=1	23.19
-----	-------

p=2	21.00
-----	-------

p=3	18.54
-----	-------

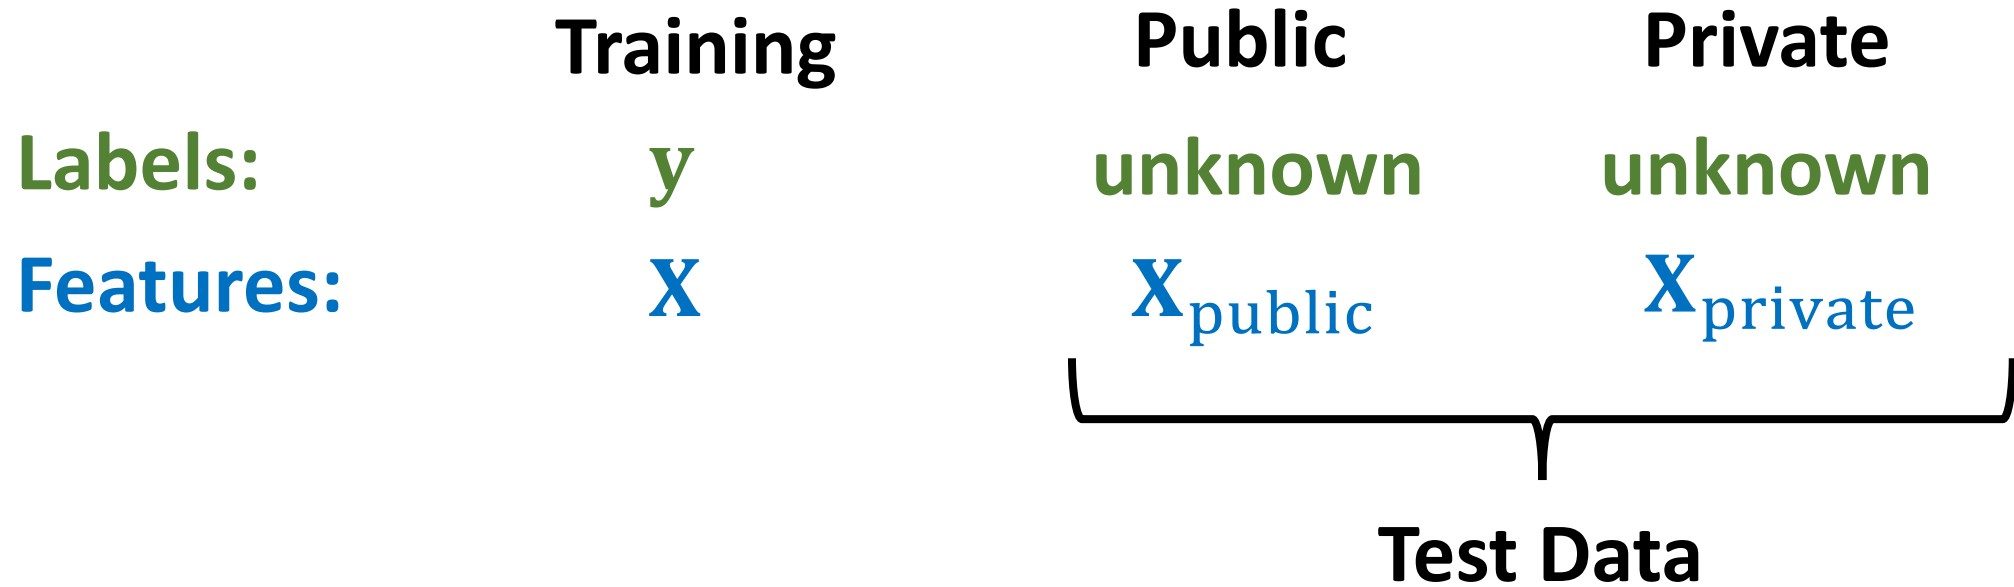
p=4	24.36
-----	-------

p=5	27.96
-----	-------

p=6	33.10
-----	-------


Real-World Machine Learning Competition

The Available Data

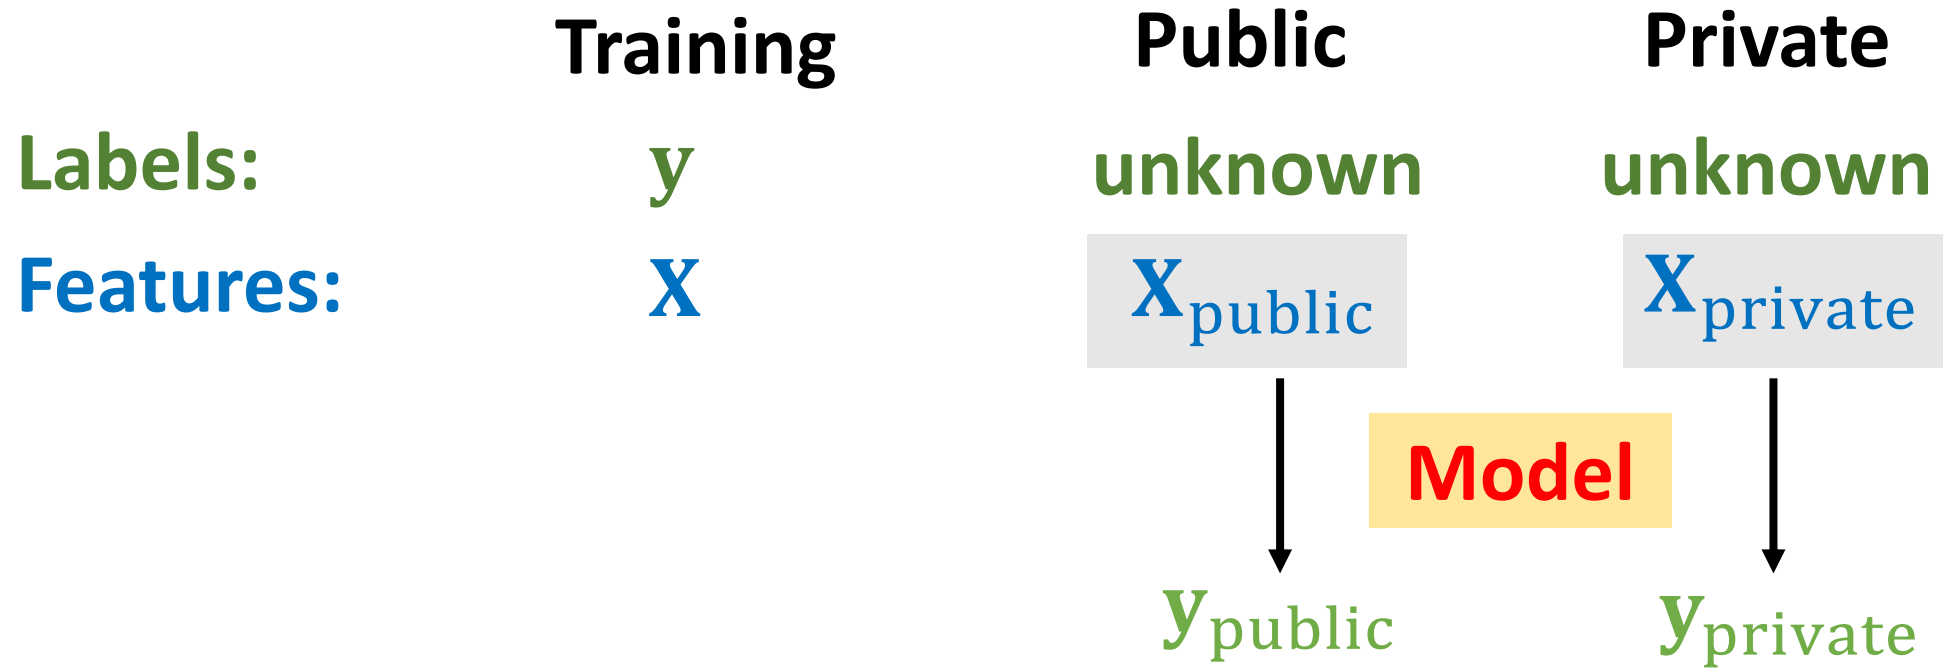


The public and private are mixed;
Participants cannot distinguish them.

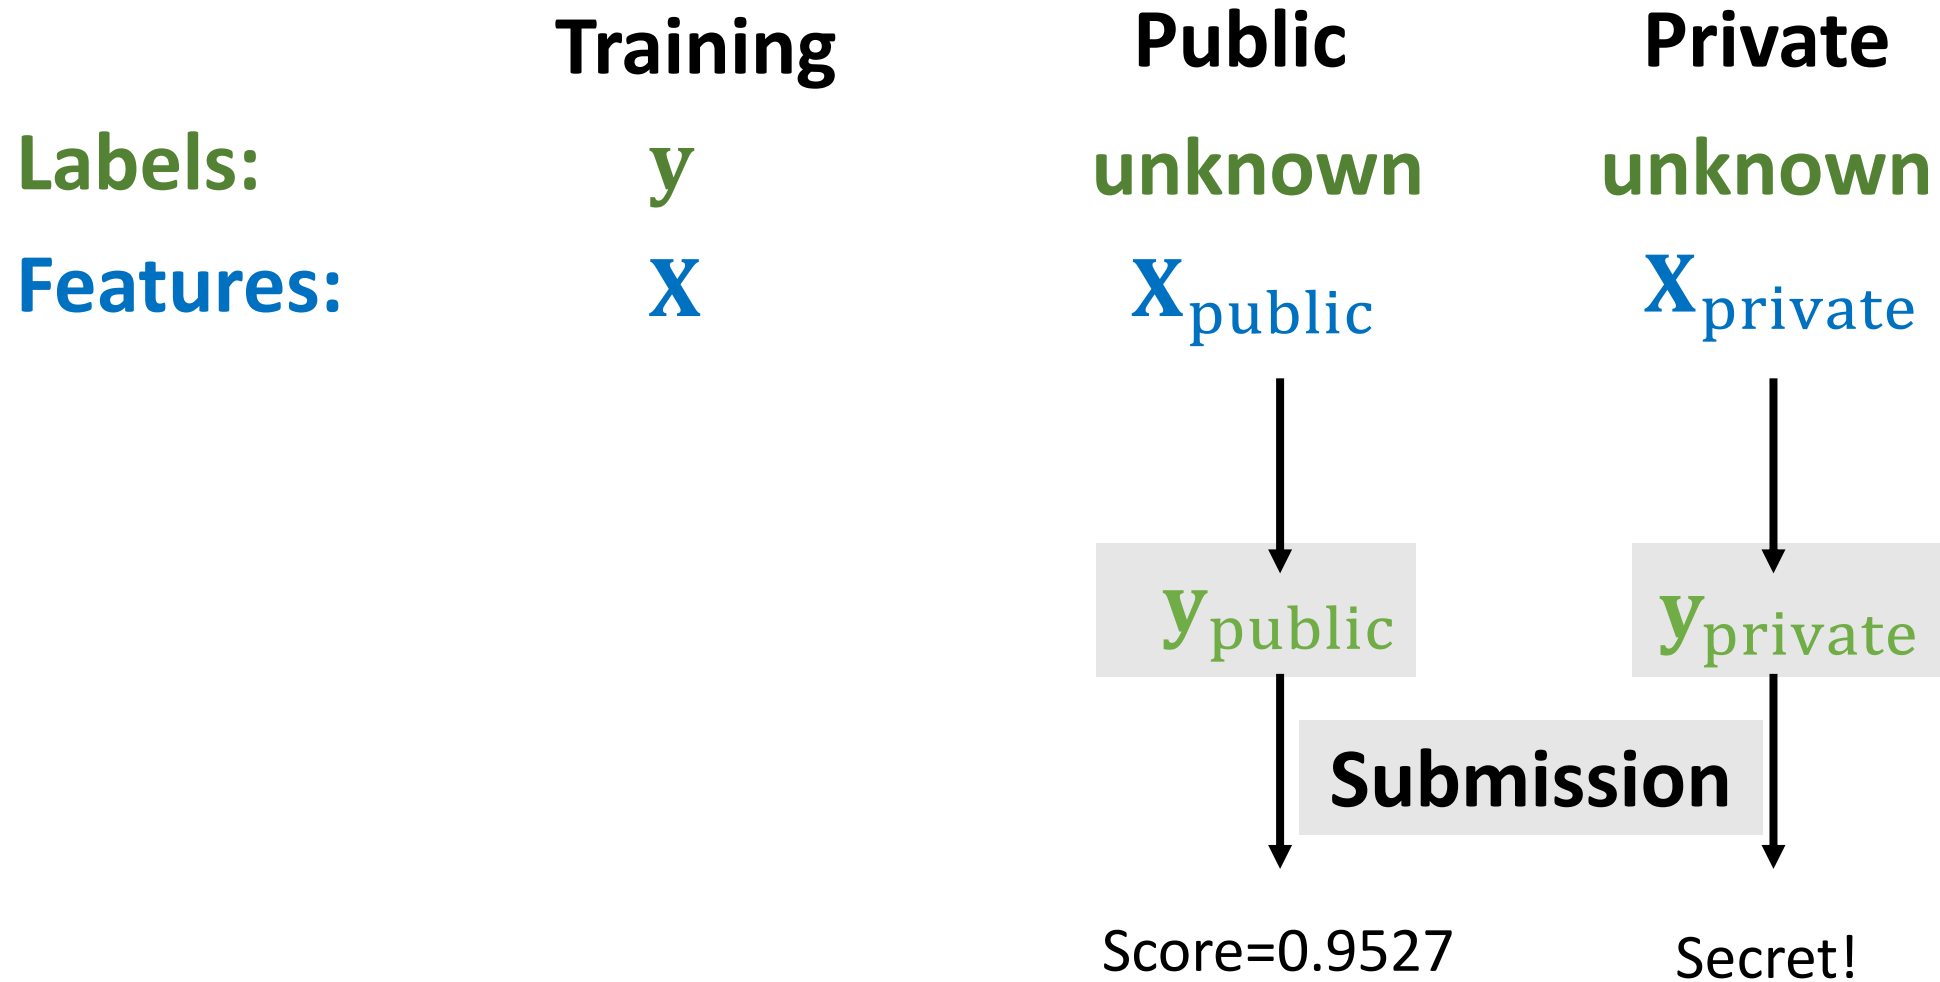
Train A Model



Prediction



Submission to Leaderboard



Submission to Leaderboard

Training

Labels:

y

Features:

X

Public

unknown

X_{public}

Private

unknown

X_{private}

Question: Why two leaderboards?

Answer: The score can be evilly used for hyper-parameter tuning (cheating).

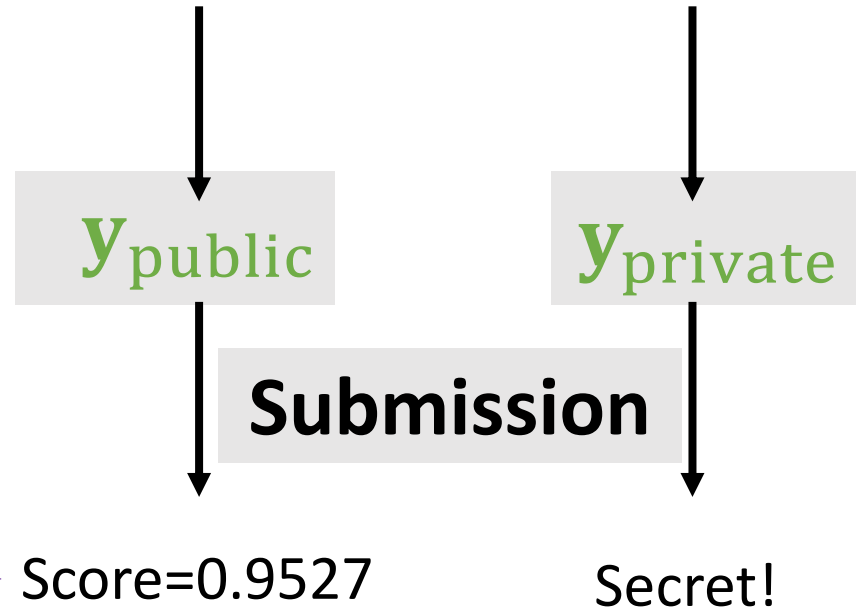
y_{public}

y_{private}

Submission

Score=0.9527

Secret!



Summary

- Polynomial regression for non-linear problems.
- Polynomial regression has a **hyper-parameter p** .
- Underfitting (very small p) and overfitting (very big p) .
- Tune the **hyper-parameters** using cross-validation.
- Make your model **parameters** and **hyper-parameters** independent of the test set!!!

Binary Classification

(Logistic Regression)

Vector and Matrix Derivatives

Derivative of Scalar w.r.t. Scalar

Examples:

- $y = x^2; \frac{dy}{dx} = 2x.$

- $y = e^x; \frac{dy}{dx} = e^x.$

Derivative of Vector w.r.t. Scalar

- The derivative of a vector $\mathbf{y} \in \mathbb{R}^n$ w.r.t. a scalar $x \in \mathbb{R}$:

$$\frac{\partial \mathbf{y}}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x} \\ \frac{\partial y_2}{\partial x} \\ \vdots \\ \frac{\partial y_n}{\partial x} \end{bmatrix}$$

- Example:

$$\mathbf{y} = \begin{bmatrix} 3x^2 \\ x + 1 \\ \log x \\ e^x \end{bmatrix}, \quad \frac{\partial \mathbf{y}}{\partial x} = \begin{bmatrix} 6x \\ 1 \\ 1/x \\ e^x \end{bmatrix}$$

Derivative of Scalar w.r.t. Vector

- The derivative of a scalar $y \in \mathbb{R}$ w.r.t. a vector $\mathbf{x} \in \mathbb{R}^m$:

$$\frac{\partial y}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_m} \end{bmatrix}$$

- Example 1:

$$y = \|\mathbf{x}\|_2^2 = \sum_{i=1}^m x_i^2, \quad \frac{\partial y}{\partial \mathbf{x}} = 2\mathbf{x}.$$

Derivative of Scalar w.r.t. Vector

- The derivative of a scalar $y \in \mathbb{R}$ w.r.t. a vector $\mathbf{x} \in \mathbb{R}^m$:

$$\frac{\partial y}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_m} \end{bmatrix}$$

- Example 2:

$$y = \mathbf{x}^T \mathbf{z} = \sum_{i=1}^m x_i z_i, \quad \frac{\partial y}{\partial \mathbf{x}} = \mathbf{z}.$$

Derivative of Scalar w.r.t. Vector

- The derivative of a scalar $y \in \mathbb{R}$ w.r.t. a vector $\mathbf{x} \in \mathbb{R}^m$:

$$\frac{\partial y}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_m} \end{bmatrix}$$

- Example 3:

$$y = \sum_{i=1}^m \log(1 + e^{-x_i}), \quad \frac{\partial y}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial \log(1+e^{-x_1})}{\partial x_1} \\ \vdots \\ \frac{\partial \log(1+e^{-x_m})}{\partial x_m} \end{bmatrix} = \begin{bmatrix} -\frac{1}{1+e^{x_1}} \\ \vdots \\ -\frac{1}{1+e^{x_m}} \end{bmatrix}$$

Derivative of Vector w.r.t. Vector

- The derivative of a vector $\mathbf{y} \in \mathbb{R}^n$ w.r.t. a vector $\mathbf{x} \in \mathbb{R}^m$:

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_2}{\partial x_1} & \cdots & \frac{\partial y_n}{\partial x_1} \\ \frac{\partial y_1}{\partial x_2} & \frac{\partial y_2}{\partial x_2} & \cdots & \frac{\partial y_n}{\partial x_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_1}{\partial x_m} & \frac{\partial y_2}{\partial x_m} & \cdots & \frac{\partial y_n}{\partial x_m} \end{bmatrix}$$

$m \times n$ matrix

- Example 1:

$$\frac{\partial \mathbf{x}}{\partial \mathbf{x}} = \underbrace{\begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix}}_{m \times m}$$

The (i, j) -th entry is $\frac{\partial y_j}{\partial x_i}$

Derivative of Vector w.r.t. Vector

- The derivative of a vector $\mathbf{y} \in \mathbb{R}^n$ w.r.t. a vector $\mathbf{x} \in \mathbb{R}^m$:

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \cdots & \frac{\partial y_1}{\partial x_m} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \cdots & \frac{\partial y_2}{\partial x_m} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_n}{\partial x_1} & \frac{\partial y_n}{\partial x_2} & \cdots & \frac{\partial y_n}{\partial x_m} \end{bmatrix} \quad m \times n \text{ matrix}$$

- Example 2:

$$\mathbf{y} = \begin{bmatrix} a_1 x_1^2 \\ a_2 x_2^2 \\ \vdots \\ a_m x_m^2 \end{bmatrix} \in \mathbb{R}^m, \quad \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \underbrace{\begin{bmatrix} 2a_1 x_1 & 0 & \cdots & 0 \\ 0 & 2a_2 x_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 2a_m x_m \end{bmatrix}}_{m \times m}$$

Derivative of Vector w.r.t. Vector

- The derivative of a vector $\mathbf{y} \in \mathbb{R}^n$ w.r.t. a vector $\mathbf{x} \in \mathbb{R}^m$:


$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \cdots & \frac{\partial y_1}{\partial x_m} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \cdots & \frac{\partial y_2}{\partial x_m} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_n}{\partial x_1} & \frac{\partial y_n}{\partial x_2} & \cdots & \frac{\partial y_n}{\partial x_m} \end{bmatrix} \quad m \times n \text{ matrix}$$

- Example 3:

$$\mathbf{A} \in \mathbb{R}^{n \times m}, \quad \mathbf{y} = \mathbf{A}\mathbf{x} \in \mathbb{R}^n, \quad \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \mathbf{A}^T \in \mathbb{R}^{m \times n}$$

Chain Rule

- Let $\mathbf{z} \in \mathbb{R}^{n_z}$ be a function of $\mathbf{y} \in \mathbb{R}^{n_y}$ and \mathbf{y} be a function of $\mathbf{x} \in \mathbb{R}^{n_x}$.



$$\underbrace{\frac{d\mathbf{z}}{d\mathbf{x}}}_{n_x \times n_z} = \underbrace{\frac{d\mathbf{y}}{d\mathbf{x}}}_{n_x \times n_y} \underbrace{\frac{d\mathbf{z}}{d\mathbf{y}}}_{n_y \times n_z}$$

Derivative of Scalar w.r.t. Matrix

- The derivative of a scalar $y \in \mathbb{R}$ w.r.t. a matrix $\mathbf{Z} \in \mathbb{R}^{p \times q}$:
 1. Vectorization: $\mathbf{x} = \text{vec}(\mathbf{Z}) \in \mathbb{R}^{pq \times 1}$.
 2. Compute $\frac{\partial y}{\partial \mathbf{x}} \in \mathbb{R}^{pq \times 1}$.
 3. Reshape the resulting $pq \times 1$ vector to $p \times q$ matrix.

Derivative of Vector w.r.t. Matrix

- The derivative of a vector $\mathbf{y} \in \mathbb{R}^n$ w.r.t. a matrix $\mathbf{Z} \in \mathbb{R}^{p \times q}$:
 1. Vectorization: $\mathbf{x} = \text{vec}(\mathbf{Z}) \in \mathbb{R}^{pq \times 1}$.
 2. Compute $\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \in \mathbb{R}^{pq \times n}$.
 3. Reshape the resulting $pq \times n$ matrix to $p \times q \times n$ tensor.

Binary Classification

Tasks

Methods

Algorithms

Binary Classification


Input: feature vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \{-1, +1\}$.

Output: a function $f: \mathbb{R}^d \mapsto \{-1, +1\}$.


Binary Classification

Input: feature vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \{-1, +1\}$.

Output: a function $f: \mathbb{R}^d \mapsto \{-1, +1\}$.



Linear Classification



Nonlinear Classification


Logistic Regression (Linear Classifier)

Tasks

Methods

Algorithms

Linear Classifier




Hyperplane: $\mathbf{w}^T \mathbf{x} + b = 0$

Define $\bar{\mathbf{x}}_j = [\mathbf{x}_j; 1] \in \mathbb{R}^{d+1}$


Define $\bar{\mathbf{w}} = [\mathbf{w}; b] \in \mathbb{R}^{d+1}$

$\Rightarrow \mathbf{x}_j^T \mathbf{w} + b = \bar{\mathbf{x}}_j^T \bar{\mathbf{w}}$


Linear Classifier




Linear Classifier



- Learn a vector \mathbf{w} such that
- If $y_j = +1$, then $\mathbf{w}^T \mathbf{x}_j > 0$.



Linear Classifier




Learn a vector \mathbf{w} such that

- If $y_j = +1$, then $\mathbf{w}^T \mathbf{x}_j > 0$.
- If $y_j = -1$, then $\mathbf{w}^T \mathbf{x}_j < 0$.

\mathbf{x}_j

\mathbf{w}

Linear Classifier



Learn a vector w such that

- If $y_j = +1$, then $w^T x_j > 0$.
- If $y_j = -1$, then $w^T x_j < 0$.




Key Idea:

Encourage $y_j w^T x_j$ to be positive

Directly Minimize the Classification Error?

Minimize $\sum_j h(y_j \mathbf{w}^T \mathbf{x}_j)$, where $h(z) = \begin{cases} 1, & \text{if } z < 0; \\ 0, & \text{if } z \geq 0. \end{cases}$




Key Idea:

Encourage $y_j \mathbf{w}^T \mathbf{x}_j$ to be positive

Directly Minimize the Classification Error?

Minimize $\sum_j h(y_j \mathbf{w}^T \mathbf{x}_j)$, where $h(z) = \begin{cases} 1, & \text{if } z < 0; \\ 0, & \text{if } z \geq 0. \end{cases}$




The objective function = # of misclassified training samples

Key Idea:

Encourage $y_j \mathbf{w}^T \mathbf{x}_j$ to be positive

Directly Minimize the Classification Error?

Minimize $\sum_j h(y_j \mathbf{w}^T \mathbf{x}_j)$, where $h(z) = \begin{cases} 1, & \text{if } z < 0; \\ 0, & \text{if } z \geq 0. \end{cases}$




The function is nonsmooth, nonconvex, and hard to optimize.

Key Idea:

Encourage $y_j \mathbf{w}^T \mathbf{x}_j$ to be positive

Logistic Regression

Minimize $\sum_j l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.



Key Idea:
Encourage $y_j \mathbf{w}^T \mathbf{x}_j$ to be positive

Logistic Regression

Tasks

Methods

Algorithms

Logistic Regression

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.

Tasks

Binary Classification

Multi-Class Classification

Methods

Logistic Regression

SVM

Neural Networks

Algorithms


Gradient Descent (GD)

Accelerated GD

Stochastic GD


Gradient

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.



Gradient


Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.



- $\frac{\partial z}{\partial \mathbf{w}} = y \mathbf{x},$

Gradient


Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.



- $\frac{\partial z}{\partial \mathbf{w}} = y \mathbf{x}, \quad \frac{\partial l(z)}{\partial z} = \frac{-e^{-z}}{1+e^{-z}} = -\frac{1}{1+e^z}.$

Gradient

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.



- $\frac{\partial z}{\partial \mathbf{w}} = y\mathbf{x}, \quad \frac{\partial l(z)}{\partial z} = \frac{-e^{-z}}{1+e^{-z}} = -\frac{1}{1+e^z}.$
- Chain rule: $\frac{\partial l}{\partial \mathbf{w}} = \frac{\partial z}{\partial \mathbf{w}} \cdot \frac{\partial l}{\partial z} = (y\mathbf{x}) \left(-\frac{1}{1+e^z} \right) = -\frac{y\mathbf{x}}{1+\exp(y\mathbf{w}^T \mathbf{x})}.$

Gradient

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.

- We have shown: $\frac{\partial l(y \mathbf{w}^T \mathbf{x})}{\partial \mathbf{w}} = \frac{-y \mathbf{x}}{1 + \exp(y \mathbf{w}^T \mathbf{x})}$.
- Objective function: $f(\mathbf{w}) = \frac{1}{n} \sum_j l(y_j \mathbf{w}^T \mathbf{x}_j)$.

Gradient

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.

- We have shown: $\frac{\partial l(y \mathbf{w}^T \mathbf{x})}{\partial \mathbf{w}} = \frac{-y \mathbf{x}}{1 + \exp(y \mathbf{w}^T \mathbf{x})}$.
- Objective function: $f(\mathbf{w}) = \frac{1}{n} \sum_j l(y_j \mathbf{w}^T \mathbf{x}_j)$.
- $\frac{\partial f(\mathbf{w})}{\partial \mathbf{w}} = \frac{1}{n} \sum_j \frac{\partial l(y_j \mathbf{w}^T \mathbf{x}_j)}{\partial \mathbf{w}} = \frac{1}{n} \sum_j \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}^T \mathbf{x}_j)}$.

Gradient

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

- We have shown: $\frac{\partial l(y \mathbf{w}^T \mathbf{x})}{\partial \mathbf{w}} = \frac{-y \mathbf{x}}{1 + \exp(y \mathbf{w}^T \mathbf{x})}$.
- Objective function: $f(\mathbf{w}) = \frac{1}{n} \sum_j l(y_j \mathbf{w}^T \mathbf{x}_j)$.
- $\frac{\partial f(\mathbf{w})}{\partial \mathbf{w}} = \frac{1}{n} \sum_j \frac{\partial l(y_j \mathbf{w}^T \mathbf{x}_j)}{\partial \mathbf{w}} = \frac{1}{n} \sum_j \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}^T \mathbf{x}_j)}$.

Gradient Descent (GD) Algorithm

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

GD repeat:

1. Compute gradient: \mathbf{g}_t
2. Update: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \mathbf{g}_t$



Tune the step size (learning rate) α

Algorithms

Gradient Descent (GD)

Accelerated GD

Stochastic GD

AGD Algorithm

Logistic regression: $\min_{\mathbf{w}} \frac{1}{n} \sum_{j=1}^n l(y_j \mathbf{w}^T \mathbf{x}_j)$, where $l(z) = \log(1 + e^{-z})$.

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

AGD repeat:

1. Compute gradient: \mathbf{g}_t
2. Update momentum: $\mathbf{v}_{t+1} = \beta \mathbf{v}_t + \mathbf{g}_t$
3. Update: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \mathbf{v}_{t+1}$

Tune α and β ($0 \leq \beta < 1$)


Algorithms

Gradient Descent (GD)

Accelerated GD

Stochastic GD

GD versus AGD



Time Complexity

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \tilde{\mathbf{g}}_{t,j}, \quad \text{where } \tilde{\mathbf{g}}_{t,j} = \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

Per-iteration time complexity is $O(nd)$.

- $O(d)$ time for computing $\mathbf{w}_t^T \mathbf{x}_j$.
- $O(d)$ time for computing $\tilde{\mathbf{g}}_{t,j}$.
- $O(nd)$ time for computing all the $\tilde{\mathbf{g}}_{t,j}$.

Algorithms

Gradient Descent (GD)

Accelerated GD

Stochastic GD

SGD Algorithm

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \tilde{\mathbf{g}}_{t,j}, \quad \text{where } \tilde{\mathbf{g}}_{t,j} = \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

The stochastic gradient is close to the full gradient:

$$\mathbf{g}_t = \mathbb{E}_j[\tilde{\mathbf{g}}_{t,j}],$$

where j is randomly sampled from $\{1, \dots, n\}$.

Algorithms

Gradient Descent (GD)

Accelerated GD

Stochastic GD

SGD Algorithm

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \tilde{\mathbf{g}}_{t,j}, \quad \text{where } \tilde{\mathbf{g}}_{t,j} = \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

SGD repeats

1. Randomly draw j from $\{1, 2, \dots, n\}$.
2. Compute the stochastic gradient $\tilde{\mathbf{g}}_{t,j}$.
3. Update: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \tilde{\mathbf{g}}_{t,j}$.

Per-iteration time complexity is $O(d)$.

Algorithms

Gradient Descent (GD)

Accelerated GD

Stochastic GD

Accelerated SGD Algorithm

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \tilde{\mathbf{g}}_{t,j}, \quad \text{where } \tilde{\mathbf{g}}_{t,j} = \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

Accelerated SGD repeats

1. Randomly draw j from $\{1, 2, \dots, n\}$.
2. Compute the stochastic gradient $\tilde{\mathbf{g}}_{t,j}$.
3. Update momentum: $\mathbf{v}_{t+1} = \beta \mathbf{v}_t + \tilde{\mathbf{g}}_{t,j}$.
4. Update: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \mathbf{v}_{t+1}$.

Algorithms

Gradient Descent (GD)

Accelerated GD

Stochastic GD

SGD Algorithm

$$\text{Gradient at } \mathbf{w}_t: \mathbf{g}_t = \frac{1}{n} \sum_{j=1}^n \tilde{\mathbf{g}}_{t,j}, \quad \text{where } \tilde{\mathbf{g}}_{t,j} = \frac{-y_j \mathbf{x}_j}{1 + \exp(y_j \mathbf{w}_t^T \mathbf{x}_j)}.$$

Output of SGD:

- Option 1: output the last iteration \mathbf{w}_{T+1}
- Option 2: output the average of \mathbf{w} produced by the last tens of iteration.


Algorithms

Gradient Descent (GD)


Accelerated GD

Stochastic GD

GD versus SGD



GD versus SGD



Training and Prediction

- Training:

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_j l(y_j \mathbf{w}^T \mathbf{x}_j), \text{ where } l(z) = \log(1 + e^{-z}).$$

- For a test feature vector $\mathbf{x}' \in \mathbb{R}^d$, make prediction by $\operatorname{sign}(\mathbf{x}'^T \mathbf{w}^*)$.

Summary

- Logistic regression model for *linear binary* classification.

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_j l(y_j \mathbf{w}^T \mathbf{x}_j), \text{ where } l(z) = \log(1 + e^{-z}).$$

- Compute the gradient using vector derivatives and the chain rule.
- Gradient-based algorithms: GD, AGD, SGD, etc.
- Make prediction using $\operatorname{sign}(\mathbf{x}'^T \mathbf{w}^*)$.

Evaluate Binary Classification

Evaluate Binary Classification


- Error Rate = $\frac{\text{\# Classification Errors}}{\text{\# Samples}}$
- Accuracy = 1 - Error Rate

Evaluate Binary Classification


- Error Rate = $\frac{\text{\# Classification Errors}}{\text{\# Samples}}$
- Accuracy = 1 - Error Rate

Error rate and **Accuracy** are not meaningful in class-imbalanced problems.

Evaluate Binary Classification



Evaluate Binary Classification



Evaluate Binary Classification

False Positive Rate (FPR)

$$= \frac{\#FP}{\#FP + \#TN}$$

True Positive Rate (TPR)


$$= \frac{\#TP}{\#TP + \#FN}$$

True Negative (TN)

False Positive (FP)

False Negative (FN)

True Positive (TP)



Evaluate Binary Classification

False Positive Rate (FPR)

$$= \frac{\#FP}{\#FP + \#TN}$$

True Negative (TN)


False Positive (FP)

True Positive Rate (TPR)

$$= \frac{\#TP}{\#TP + \#FN}$$

False Negative (FN)

True Positive (TP)



Evaluate Binary Classification

False Positive Rate (FPR)

$$= \frac{\#FP}{\#FP + \#TN}$$

Healthy (negative)

$$y = -1$$

True Positive Rate (TPR)

$$= \frac{\#TP}{\#TP + \#FN}$$

True Negative (TN)


False Positive (FP)

False Negative (FN)

True Positive (TP)

Cancer (positive)

$$y = +1$$



Evaluate Binary Classification

False Positive Rate (FPR)

$$= \frac{\#FP}{\#FP + \#TN}$$

True Positive Rate (TPR)


$$= \frac{\#TP}{\#TP + \#FN}$$

True Negative (TN)

False Positive (FP)

False Negative (FN)

True Positive (TP)



Evaluate Binary Classification

False Positive Rate (FPR)

$$= \frac{\#FP}{\#FP + \#TN}$$

True Positive Rate (TPR)


$$= \frac{\#TP}{\#TP + \#FN}$$

True Negative (TN)


False Positive (FP)

False Negative (FN)


True Positive (TP)




Receiver Operating Characteristic (ROC) Curve




Receiver Operating Characteristic (ROC) Curve




Receiver Operating Characteristic (ROC) Curve




Receiver Operating Characteristic (ROC) Curve



Receiver Operating Characteristic (ROC) Curve



Receiver Operating Characteristic (ROC) Curve



Thank you!